

Crunching and visualizing Big Data on a Computer Cluster

Joana Simões

April 10, 2015

https://github.com/doublebyte1/workshop_bdigital



Table of Contents

- 1 Introduction
- 2 Importing a Spatial-Temporal Series
- 3 Recovering the Spatial Attributes
- 4 Putting it All Together
- 5 Piping the Results into the Outside World

Motivation

- Problem: the increasing volume of information, by explosion of traditional sources + new sources
- Target: fast query responses, which require a scalable architecture
- Possible solution: support clusters on a cost-effective architecture, such as commodity clusters or cloud environments



Cloud Services



A thought...

First the use case, then the tools.

Use Case

- Study spatial and temporal patterns of road traffic accidents.
- Relate target variable (accident) with context variables (e.g.: weather, proximity to SPI).

Use Case

- Study spatial and temporal patterns of road traffic accidents.
- Relate target variable (accident) with context variables (e.g.: weather, proximity to SPI).
- In most (Big) Data Analysis 80% of the effort is devoted to the Extract-Transform-Load (ETL) process
- ETL: process responsible for pulling data out of the source systems and placing it into a data warehouse

Use Case

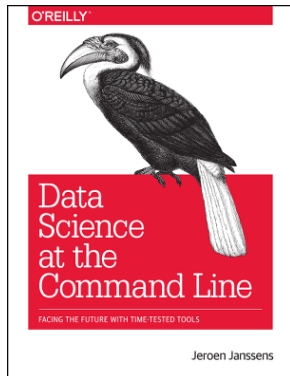
- Study spatial and temporal patterns of road traffic accidents.
- Relate target variable (accident) with context variables (e.g.: weather, proximity to SPI).
- In most (Big) Data Analysis 80% of the effort is devoted to the Extract-Transform-Load (ETL) process
- ETL: process responsible for pulling data out of the source systems and placing it into a data warehouse
 - **Extract** data from different source systems and convert it into one consolidated data warehouse format which is ready for transformation processing
 - **Transform**: cleaning, filtering, splitting a column, joining data, apply validation, apply rules, etc
 - **Load**: into the data warehouse, repository or reporting applications

Another thought...

There are no free lunches.

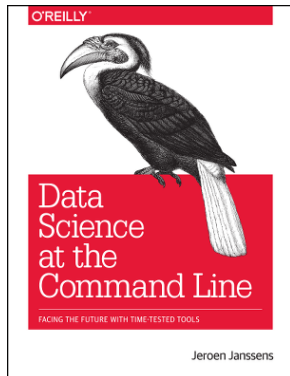
Scalable ML Platforms

- ML platforms may provide high-level data import tools:



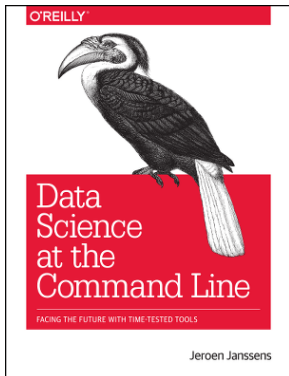
Scalable ML Platforms

- ML platforms may provide high-level data import tools:
 - They generally trade ease of use for flexibility




Scalable ML Platforms

- ML platforms may provide high-level data import tools:
 - They generally trade ease of use for flexibility
 - They may have a cost (\$\$\$)

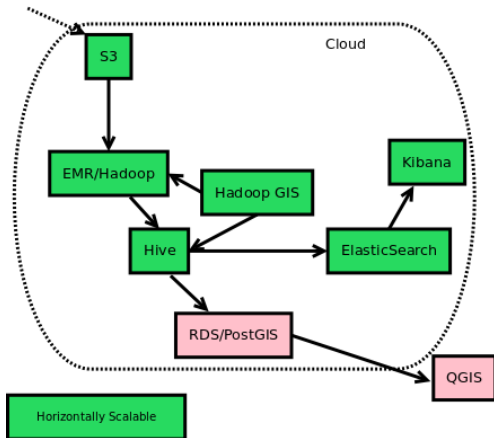


Scalable ML Platforms

- ML platforms may provide high-level data import tools:
 - They generally trade ease of use for flexibility
 - They may have a cost (\$\$\$)
- To ensure maximum flexibility, we should be able to link together many tools, often using the command line.

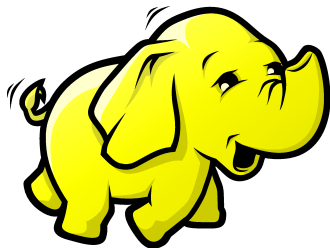


Stack



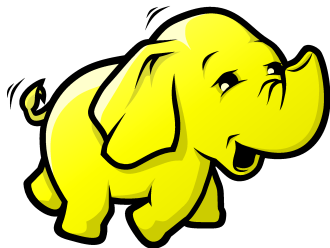
Apache Hadoop

- framework for distributed storage and processing of (Big) Data on computer Clusters



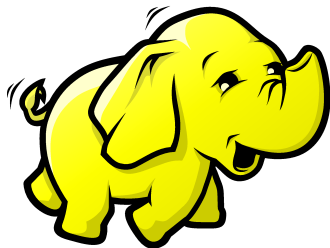
Apache Hadoop

- framework for distributed storage and processing of (Big) Data on computer Clusters
 - **Storage:** HDFS
 - **Processing:** MapReduce



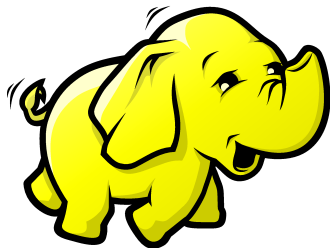
Apache Hadoop

- framework for distributed storage and processing of (Big) Data on computer Clusters
 - **Storage:** HDFS
 - **Processing:** MapReduce
- It features a FOS license (Apache 2.0)



Apache Hadoop

- framework for distributed storage and processing of (Big) Data on computer Clusters
 - **Storage:** HDFS
 - **Processing:** MapReduce
- It features a FOS license (Apache 2.0)
- EMR is an Amazon service that uses Hadoop



Apache Hive

- Data warehouse infrastructure built on top of Hadoop for providing data summarization, query, and analysis.



Apache Hive

- Data warehouse infrastructure built on top of Hadoop for providing data summarization, query, and analysis.



Apache Hive

- Data warehouse infrastructure built on top of Hadoop for providing data summarization, query, and analysis.
- HiveQL: query language based on SQL
 - An internal compiler translates HiveQL statements into MapReduce jobs
- It features a FOS license (Apache 2.0)



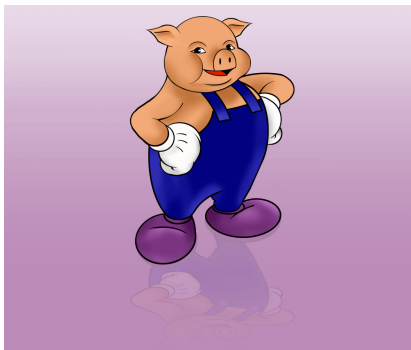
Apache Hive

- Data warehouse infrastructure built on top of Hadoop for providing data summarization, query, and analysis.
- HiveQL: query language based on SQL
 - An internal compiler translates HiveQL statements into MapReduce jobs
- It features a FOS license (Apache 2.0)
- EMR features an Hive installation



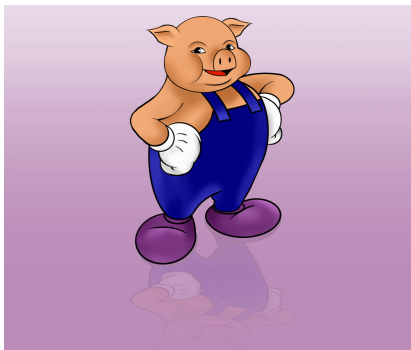
Apache Pig

- High-level platform for creating MapReduce jobs over Hadoop.



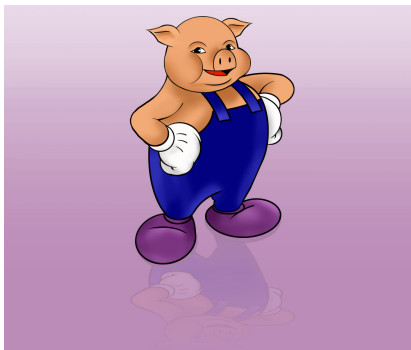
Apache Pig

- High-level platform for creating MapReduce jobs over Hadoop.
- It features an interactive mode and a batch mode



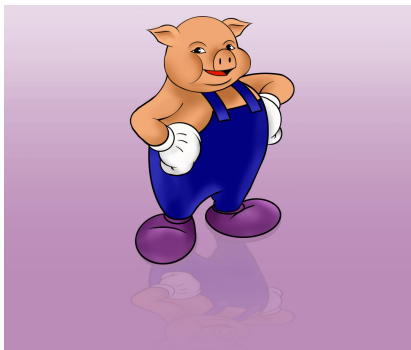
Apache Pig

- High-level platform for creating MapReduce jobs over Hadoop.
- It features an interactive mode and a batch mode
- It uses lazy evaluation



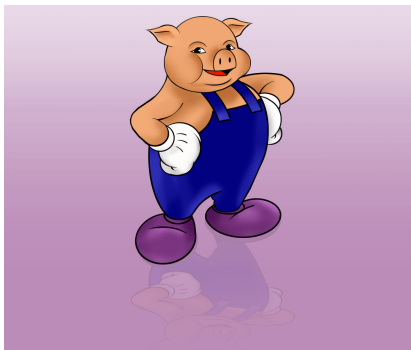
Apache Pig

- High-level platform for creating MapReduce jobs over Hadoop.
- It features an interactive mode and a batch mode
- It uses lazy evaluation
- Pig Latin is procedural



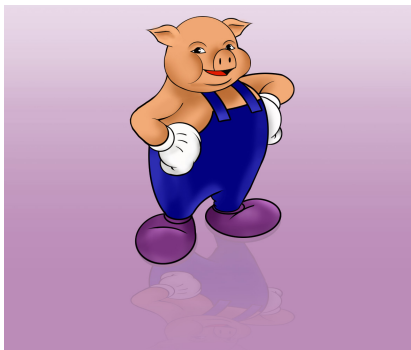
Apache Pig

- High-level platform for creating MapReduce jobs over Hadoop.
- It features an interactive mode and a batch mode
- It uses lazy evaluation
- Pig Latin is procedural
- It features a FOS license (Apache 2.0)



Apache Pig

- High-level platform for creating MapReduce jobs over Hadoop.
- It features an interactive mode and a batch mode
- It uses lazy evaluation
- Pig Latin is procedural
- It features a FOS license (Apache 2.0)
- EMR may install Pig



Spatial Support



Hadoop GIS

- FOS toolkit for "Big Spatial Data Analytics".



Hadoop GIS

- FOS toolkit for "Big Spatial Data Analytics".
- Hosted on github.



Hadoop GIS

- FOS toolkit for "Big Spatial Data Analytics".
- Hosted on github.
- It consists in three libraries:



Hadoop GIS

- FOS toolkit for "Big Spatial Data Analytics".
- Hosted on github.
- It consists in three libraries:
 - Esri Geometry API for Java.
 - Geoprocessing Tools for Hadoop.
 - **Spatial Framework for Hadoop (SFH)**: extends Hive to enable spatial queries and geometry types.



Pigeon

- Wrapper around the geometry API from ESRI.
- Adds spatial support to Pig Latin.

Systems: Pigeon



- **Pigeon : Spatial Support in Pig**

- University of Minnesota (Prof. Mohamed Mokbel)

- **Pig Query(Non Spatial)**

```
points = LOAD 'points' AS (id,long, lon,double, lat,double);  
results = FILTER points BY  
    lon < -93.158 AND lon > -93.175 AND  
    lat > 45.0077 AND lat < 45.0164;  
STORE results INTO 'results';
```

- **Pigeon Query(Spatial)**

```
IMPORT 'pigeon_import.pig';  
points = LOAD 'points-pigeon' AS (id,long, location);  
results = FILTER points BY  
    ST_Contains(ST_MakeBox(-93.175, 45.0077, -93.158, 45.0164), location);  
STORE results INTO 'results-pigeon';
```

PostGIS

- Spatial extension for the RDBMS PostgreSQL
- Free and Open Source License (GPL 2.0)



PostGIS

- Spatial extension for the RDBMS PostgreSQL
- Free and Open Source License (GPL 2.0)
- Standards support



PostGIS

- Spatial extension for the RDBMS PostgreSQL
- Free and Open Source License (GPL 2.0)
- Standards support
- High Performance



PostGIS

- Spatial extension for the RDBMS PostgreSQL
- Free and Open Source License (GPL 2.0)
- Standards support
- High Performance
- Famous users: Foursquare, Instagram, CartoDB.



E(L)K

- **ElasticSearch:** Search engine based on Lucene (RESTful, JSON).



E(L)K

- **ElasticSearch:** Search engine based on Lucene (RESTful, JSON).
- **Logstash:** tool for managing events and logs.



E(L)K

- **ElasticSearch:** Search engine based on Lucene (RESTful, JSON).
- **Logstash:** tool for managing events and logs.
- **Kibana:** data visualization platform (custom dashboards).



Practical

Hands-on

Connect to the Cluster

- Micro-task: Connect to the cluster using SSH

Connect to the Cluster

- Micro-task: Connect to the cluster using SSH
 - Start the SSH client (differences may apply), and connect using the pem certificate.
 - <https://drive.google.com/a/bdigital.org/file/d/0B5kkho5DSzlyanBkUzgtRWh20VE/view?usp=sharing>

From S3 to HDFS

- Micro-task: Understand the dataset structure
 - A sample dataset is stored on an S3 bucket:
`s3n://workshop-bdsd/accidents/`

From S3 to HDFS

- Micro-task: Understand the dataset structure
 - A sample dataset is stored on an S3 bucket:
`s3n://workshop-bdsd/accidents/`
 - `https://s3-eu-west-1.amazonaws.com/workshop-bdsd/accidents/accidents_sample.csv`

From S3 to HDFS

- Micro-task: Understand the dataset structure
 - A sample dataset is stored on an S3 bucket:
`s3n://workshop-bdsd/accidents/`
 - `https://s3-eu-west-1.amazonaws.com/workshop-bdsd/accidents/accidents_sample.csv`
 - Download and view dataset

From S3 to HDFS (cont.)

- Micro-task: Create a table linking to the data
 - Enter hive and create an external table linking to the S3 bucket

From S3 to HDFS (cont.)

- Micro-task: Create a table linking to the data
 - Enter hive and create an external table linking to the S3 bucket
 - Use the CSV serde to parse the table structure

From S3 to HDFS (cont.)

- Micro-task: Create a table linking to the data
 - Enter hive and create an external table linking to the S3 bucket
 - Use the CSV serde to parse the table structure
 - separator char
 - quote char
 - headers

From S3 to HDFS (cont.)

- Micro-task: Create a table linking to the data
 - Enter hive and create an external table linking to the S3 bucket
 - Use the CSV serde to parse the table structure
 - separator char
 - quote char
 - headers
 - View imported data

From S3 to HDFS (cont.)

- Micro-task: Type Mapping

From S3 to HDFS (cont.)

- Micro-task: Type Mapping
 - Create an empty table with correct types

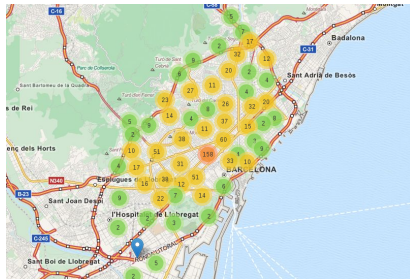
From S3 to HDFS (cont.)

- Micro-task: Type Mapping
 - Create an empty table with correct types
 - Insert data from `accidents_import`

From S3 to HDFS (cont.)

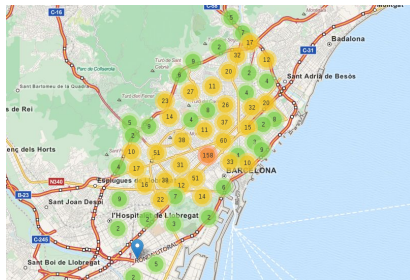
- Micro-task: Type Mapping
 - Create an empty table with correct types
 - Insert data from accidents_import
 - View table

What is so "Special" about Spatial



What is so "Special" about Spatial

- Location attributes allow us to detect spatial patterns
- Location also works as a "key", allowing us to connect with other datasets



Analysis of the Spatial Attributes

- Spatial Attributes are encoded as coordinates in "d_coord_geo_impacte"

Analysis of the Spatial Attributes

- Spatial Attributes are encoded as coordinates in "d_coord_geo_impacte"
- Problems with this field:

Analysis of the Spatial Attributes

- Spatial Attributes are encoded as coordinates in "d_coord_geo_impacte"
- Problems with this field:
 - Unstructured: needs parsing
 - Inconsistent format (order of coordinates, separator)
 - Invalid values (e.g.: 77, names)
 - No metadata

Analysis of the Spatial Attributes

- Spatial Attributes are encoded as coordinates in "d_coord_geo_impacte"
- Problems with this field:
 - Unstructured: needs parsing
 - Inconsistent format (order of coordinates, separator)
 - Invalid values (e.g.: 77, names)
 - No metadata
 - Mixed CRS:

Analysis of the Spatial Attributes

- Spatial Attributes are encoded as coordinates in "d_coord_geo_impacte"
- Problems with this field:
 - Unstructured: needs parsing
 - Inconsistent format (order of coordinates, separator)
 - Invalid values (e.g.: 77, names)
 - No metadata
 - Mixed CRS:
 - WGS84 (EPSG:4326)
 - European Grid (EPSG:5554)
 - European Grid encoded by the police using an *ad-hoc* format


Analysis of the Spatial Attributes

- Spatial Attributes are encoded as coordinates in "d_coord_geo_impacte"
- Problems with this field:
 - Unstructured: needs parsing
 - Inconsistent format (order of coordinates, separator)
 - Invalid values (e.g.: 77, names)
 - No metadata
 - Mixed CRS:
 - WGS84 (EPSG:4326)
 - European Grid (EPSG:5554)
 - European Grid encoded by the police using an *ad-hoc* format
 - $lon = y/1000 + 400000$
 - $lat = y/1000 + 4500000$

CRS

World Geodetic System (WGS84, EPSG:4326): standard for use in cartography, geodesy, and navigation; reference CRS for GPS.

European Terrestrial Reference System 1989 (ETRS89, EPSG:5554): proposed, multipurpose Pan-European mapping standard; based on the ETRS89 Lambert Azimuthal Equal-Area projection coordinate reference system



Objective

- Separate lat, long fields and map them to correct types
- Remove invalid values
- Convert all coordinates into a single CRS (WGS84)

Exporting Data to Pig

- Pig uses filters to subset the data
- To merge back the subsetting data, we can use joins by a common field
- Micro-task: Export the data

Exporting Data to Pig

- Pig uses filters to subset the data
- To merge back the subsetting data, we can use joins by a common field
- Micro-task: Export the data
 - Create a copy of the accidents table, with an id field (joins).

Exporting Data to Pig


- Pig uses filters to subset the data
- To merge back the subsetting data, we can use joins by a common field
- Micro-task: Export the data
 - Create a copy of the accidents table, with an id field (joins).
 - Export this table into a tsv

Exporting Data to Pig

- Pig uses filters to subset the data
- To merge back the subsetting data, we can use joins by a common field
- Micro-task: Export the data
 - Create a copy of the accidents table, with an id field (joins).
 - Export this table into a tsv
 - Store it in HDFS (if needed)
 - View exported data

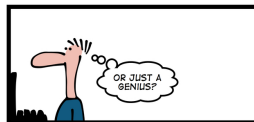

Presenting the Pig Script

- Subsets the coordinate list, using filters
- Detects each coordinate "type", using regular expressions
- In the case of grid encoded, it applies a formula to decode back into grid
- Stores the results into separate files, in HDFS



REGEX



- Sequence of characters that forms a search pattern, mainly for use in pattern matching with strings, or string matching



YESTERDAYS REGEX

REGEX



- Sequence of characters that forms a search pattern, mainly for use in pattern matching with strings, or string matching
- `REGEX_EXTRACT(D_COORD_(z]',0)`



YESTERDAYS REGEX

REGEX

- Sequence of characters that forms a search pattern, mainly for use in pattern matching with strings, or string matching
- `REGEX_EXTRACT(D_COORD_Cz]',0)`
- `'[A-z]'`



YESTERDAYS REGEX

Running Pig

- Micro-task: run pig script

Running Pig

- Micro-task: run pig script
 - Download script from S3: https://s3-eu-west-1.amazonaws.com/workshop-bdsd/recover_geography.pig

Running Pig

- Micro-task: run pig script
 - Download script from S3: `https://s3-eu-west-1.amazonaws.com/workshop-bdsd/recover_geography.pig`
 - Edit script and **ammend paths**

Running Pig

- Micro-task: run pig script
 - Download script from S3: `https://s3-eu-west-1.amazonaws.com/workshop-bdsd/recover_geography.pig`
 - Edit script and **ammend paths**
 - Run script

Running Pig

- Micro-task: run pig script
 - Download script from S3: `https://s3-eu-west-1.amazonaws.com/workshop-bdsd/recover_geography.pig`
 - Edit script and **ammend paths**
 - Run script
 - Check output files

Importing Data Back into Hive

- Micro-task: Create tables linking to pig output

Importing Data Back into Hive

- Micro-task: Create tables linking to pig output
 - Create table with wgs84 data
 - Create table with grid data
 - Create table with police-decoded data

Exporting data into PostGIS

- As of Hadoop GIS 2.0, CRS transformation is **not supported**
- We need to rely on another tool: PostGIS on RDS
- Micro-task: Export grid data to PostGIS



Exporting data into PostGIS

- As of Hadoop GIS 2.0, CRS transformation is **not supported**
- We need to rely on another tool: PostGIS on RDS
- Micro-task: Export grid data to PostGIS
 - Merge grid (grid + police) tables in a single table



Exporting data into PostGIS

- As of Hadoop GIS 2.0, CRS transformation is **not supported**
- We need to rely on another tool: PostGIS on RDS
- Micro-task: Export grid data to PostGIS
 - Merge grid (grid + police) tables in a single table
 - Exported merged table into TSV



Importing Data into PostGIS

- Micro-task: Import grid data into PostGIS

Importing Data into PostGIS

- Micro-task: Import grid data into PostGIS
 - Install the PSQL client
 - Log into RDS:

Importing Data into PostGIS

- Micro-task: Import grid data into PostGIS
 - Install the PSQL client
 - Log into RDS:
 - host: `bdigitaldb.celqzuwfokoe.eu-west-1.rds.amazonaws.com`
 - user: `workshop`
 - password: `geohipster`
 - database: `workshop_bdigital`

Importing Data into PostGIS

- Micro-task: Import grid data into PostGIS
 - Install the PSQL client
 - Log into RDS:
 - host: `bdigitaldb.celqzuwfokoe.eu-west-1.rds.amazonaws.com`
 - user: `workshop`
 - password: `geohipster`
 - database: `workshop_bdigital`
 - Create table to accomodate data

Importing Data into PostGIS

- Micro-task: Import grid data into PostGIS
 - Install the PSQL client
 - Log into RDS:
 - host: `bdigitaldb.celqzuwfokoe.eu-west-1.rds.amazonaws.com`
 - user: `workshop`
 - password: `geohipster`
 - database: `workshop_bdigital`
 - Create table to accomodate data
 - Copy data into table

CRS Transformation

- Micro-task: Convert all features in European grid to WGS84

CRS Transformation

- Micro-task: Convert all features in European grid to WGS84
 - Create geometry fields to accomodate geometry in the two CRS (Grid, WGS84)

CRS Transformation

- Micro-task: Convert all features in European grid to WGS84
 - Create geometry fields to accomodate geometry in the two CRS (Grid, WGS84)
 - Add columns
 - Set SRID
 - Create geometry index

CRS Transformation

- Micro-task: Convert all features in European grid to WGS84
 - Create geometry fields to accomodate geometry in the two CRS (Grid, WGS84)
 - Add columns
 - Set SRID
 - Create geometry index
 - Instantiate grid geometry

CRS Transformation

- Micro-task: Convert all features in European grid to WGS84
 - Create geometry fields to accomodate geometry in the two CRS (Grid, WGS84)
 - Add columns
 - Set SRID
 - Create geometry index
 - Instantiate grid geometry
 - Transform grid geometry into another CRS

CRS Transformation

- Micro-task: Convert all features in European grid to WGS84
 - Create geometry fields to accomodate geometry in the two CRS (Grid, WGS84)
 - Add columns
 - Set SRID
 - Create geometry index
 - Instantiate grid geometry
 - Transform grid geometry into another CRS
 - Export grid geometry in GeoJSON

GeoJSON

GeoJSON: is an open standard format for encoding collections of simple geographical features along with their non-spatial attributes using JavaScript Object Notation.

The screenshot shows the GeoJSONLint website (geojsonlint.com). The interface includes a navigation bar with links for Point, LineString, Polygon, Feature, FeatureCollection, and GeometryCollection. Below the navigation bar, there is a text area containing a GeoJSON FeatureCollection with three features: a Point in Sydney, Australia; a Point in Cape Town, South Africa; and a LineString connecting Sydney and Cape Town. Below the text area are buttons for 'Test GeoJSON' and 'Clear', and a checkbox for 'Clear Current Features'. To the right of the text area is a world map with three blue location pins: one in Sydney, one in Cape Town, and one in Australia. A blue line connects the Sydney and Cape Town pins. The map is powered by Leaflet, with tiles courtesy of MapQuest and map data from OpenStreetMap contributors.

```
{
  "type": "Point",
  "id": "Sydney, Australia",
  "properties": {},
  "type": "Feature"
},
{
  "geometry": {
    "coordinates": [ 18.4172485,
      -33.9289049
    ],
    "type": "Point",
    "id": "Cape Town, South Africa",
    "properties": {},
    "type": "Feature"
  },
  "geometry": {
    "coordinates": [
      [151.2164539, -33.8548157],
      [18.4172485, -33.9289049]
    ],
    "type": "LineString",
    "id": null,
    "properties": {},
    "type": "FeatureCollection"
  }
}
```

Use this site to validate and view your GeoJSON. For details about GeoJSON, [read the spec.](#)

Clear Current Features

Test GeoJSON Clear

Powered by Leaflet — Tiles Courtesy of MapQuest, Map data (c) OpenStreetMap contributors, CC-BY-SA

Importing Data back into Hive

- Micro-task: Import transformed data

Importing Data back into Hive

- Micro-task: Import transformed data
 - Enter Hive

Importing Data back into Hive

- Micro-task: Import transformed data
 - Enter Hive
 - Create table linking to the PostGIS export

Importing Data back into Hive

- Micro-task: Import transformed data
 - Enter Hive
 - Create table linking to the PostGIS export
 - Create new table and instantiate geometry from GeoJSON

Joining Data

- Micro-task: Join imported coordinates with WGS84 coordinates and the rest of the dataset

Joining Data

- Micro-task: Join imported coordinates with WGS84 coordinates and the rest of the dataset
 - Join imported records with original table with all fields

Joining Data


- Micro-task: Join imported coordinates with WGS84 coordinates and the rest of the dataset
 - Join imported records with original table with all fields
 - Merge imported records with WGS84 records, for a single table with unified geometry

ElasticSearch

Indexing the results in ElasticSearch



And now Kibana...!



Thank you for Listening!

